

Artificial Intelligence 1

Prof. Dr. Frank Hopfgartner

Dr. Matthias Horbach

Institute for Web Science and Technologies (WeST)
University of Koblenz

- 1 Introduction
- 2 Classical logics and Prolog
- 3 Search and automatic planning
- 4 Knowledge representation and reasoning
 - Default logic
 - Answer set programming
 - Argumentation
- 5 Agents and multi agent systems
- 6 Summary and conclusion

- ▶ Definition: *Knowledge representation and reasoning (KR) is the field of artificial intelligence (AI) dedicated to representing information about the world in a form that a computer system can utilize to solve complex tasks such as diagnosing a medical condition.* [Wikipedia]
- ▶ Goal: Formalisation of beliefs (knowledge) and automatic reasoning
- ▶ But we already looked at propositional and first-order logic, isn't that enough?

Every logic (=formal system) has the following components:

1. Syntax: What are the possible statements?
 - 1.1 Signature: What symbols are allowed?
($S = \{\text{Anna, human, student}\}$)
 - 1.2 Grammar: how can symbols be combined in order to obtain complex statements?
($\text{student} \Rightarrow \text{human}$)
2. Semantics: Which are the “true” statements? What is the relationship between true statements?
 - 2.1 Interpretations: Which symbol represents which concrete object?
($\text{Anna} = \text{“Anna Schmidt”}$)
 - 2.2 Models: Which statement is true in a given constellation of objects?
 - 2.3 Reasoning: How can we infer new information?
3. Calculus: How can “reasoning” be implemented?

Limitations of classical logic 1/4

- ▶ Classical logic is monotonic:

$$\text{If } \alpha \vdash \beta \quad \text{then} \quad \alpha \wedge \gamma \vdash \beta$$

for every formula γ .

- ▶ In other words: inferences are never retracted when new information is received
- ▶ Classical logic is binary:

$$\alpha \wedge (\alpha \Rightarrow \beta) \vdash \beta$$

is always true without exception.

- ▶ Central concept in knowledge representation are *rules*
- ▶ Rules always have exceptions (and there are also exceptions for that)

Limitations of classical logic 2/4

Example

Let us model our knowledge about some animals:

- birds are animals
- penguins are birds
- birds usually fly
- penguins do not fly

Naive formalisation:

$$\forall X : (bird(X) \Rightarrow animal(X))$$

$$\forall X : (penguin(X) \Rightarrow bird(X))$$

$$\forall X : (bird(X) \Rightarrow flies(X))$$

$$\forall X : (penguin(X) \Rightarrow \neg flies(X))$$

$$penguin(tweety)$$

What is the problem?

Limitations of classical logic 3/4

Example

Knowledge base in propositional logic (simpler)

$$\begin{aligned}\phi = & (bird \Rightarrow animal) \wedge (penguin \Rightarrow bird) \wedge (bird \Rightarrow flies) \\ & \wedge (penguin \Rightarrow \neg flies) \wedge penguin\end{aligned}$$

Observation:

- ▶ Knowledge base is inconsistent: $\phi \vdash flies$ and $\phi \vdash \neg flies$
- ▶ The “rule” $bird \Rightarrow flies$ is not a classical implication, it is not universally valid
- ▶ What now?
- ▶ Model exceptions explicitly: $bird \Rightarrow flies \longrightarrow bird \wedge \neg penguin \Rightarrow flies$
- ▶ What about emus, dodos, ostriches, ...?

- ▶ Explicit enumeration of exceptions is not feasible (hard to maintain, ...)
- ▶ How do we humans do it?
- ▶ When encountering a new bird for the first time, we make the *default assumption* that it can fly
- ▶ Only when obtaining explicit information that the bird does not fly, we will revise our inference

- 1 Introduction
- 2 Classical logics and Prolog
- 3 Search and automatic planning
- 4 Knowledge representation and reasoning
 - Default logic
 - Answer set programming
 - Argumentation
- 5 Agents and multi agent systems
- 6 Summary and conclusion

- ▶ Default logics are logics that allow for *non-monotonic* reasoning

Even if $\alpha \vdash \beta$ it may be that $\alpha \wedge \gamma \not\vdash \beta$

- ▶ In the following, we consider the prototype of a default logic: Reiter's default logic.

R. Reiter. A logic for default reasoning. Artificial Intelligence, 13:81–132, 1980.

Some additional notation for first-order logic 1/2

Let $\Sigma = (U, P, F)$ be a first-order signature, V a set of variables, and $\mathcal{L}(\Sigma, V)$ the corresponding first-order language.

Definition

A variable $X \in V$ is *free* in formula ϕ if there is an occurrence of X in ϕ which is not quantified. A variable $X \in V$ is *bound* in ϕ if it occurs in a sub-formula of the type $\forall X : \phi'$ or $\exists X : \phi'$.

A formula $\phi \in \mathcal{L}(\Sigma, V)$ is *closed* if every appearing variable is bound.

Example

The following two formulas are closed:

$$\forall X : (a(X) \wedge \forall Y : b(X, Y)) \qquad r(s, t) \vee \exists Y : a(Y)$$

The following formulas are not closed

$$(\forall X : a(X)) \wedge b(X) \qquad r(s, Y) \vee \exists X : a(X)$$

Some additional notation for first-order logic 2/2

A set of first-order formulas is usually seen as equivalent with the conjunction of its elements

$$\{\phi_1, \phi_2, \phi_3\} \vdash \psi \quad \Longleftrightarrow \quad \phi_1 \wedge \phi_2 \wedge \phi_3 \vdash \psi$$

Definition

The *deductive closure* of a set of first-order formulas $\Psi \subseteq \mathcal{L}(\Sigma, V)$ is defined via $Cn(\Psi) = \{\phi \mid \Psi \vdash \phi\}$.

Observe that $\Psi \subseteq Cn(\Psi)$ is always the case.

Definition

A set of first-order formulas $\Psi \subseteq \mathcal{L}(\Sigma, V)$ is *deductively closed* if $Cn(\Psi) = \Psi$.

Definition

Let $\phi, \psi_1, \dots, \psi_n, \chi \in \mathcal{L}(\Sigma, V)$ be closed first-order formulas (or propositional formulas).

Then

$$\delta = \frac{\phi : \psi_1, \dots, \psi_n}{\chi}$$

is called a *default rule* (or just *default*).

Meaning: If ϕ is known and ψ_1, \dots, ψ_n can be consistently assumed, then conclude χ .

$\phi = \text{pre}(\delta)$ (default) precondition

$\chi = \text{cons}(\delta)$ (default) conclusion

$\{\psi_1, \dots, \psi_n\} = \text{just}(\delta)$ (default) justifications

Default rules - example

- ▶ Rules with exceptions:

$$\frac{\textit{bird} : \textit{flies}}{\textit{flies}}$$

- ▶ Rules that are typically true:

$$\frac{\textit{going_to_work} : \textit{take_bus}}{\textit{take_bus}}$$

- ▶ Rules that hold unless the opposite can be shown:

$$\frac{\textit{accused} : \textit{innocent}}{\textit{innocent}}$$

A default

$$\delta = \frac{\phi : \psi_1, \dots, \psi_n}{\chi}$$

with *open* formulas $\phi, \psi_1, \dots, \psi_n$, χ is interpreted as a schema, i. e., as its set of grounded defaults (over the given universe).

Default schema - example

The default rule

$$\frac{\text{friend}(X, Y) \wedge \text{friend}(Y, Z) : \text{friend}(X, Z)}{\text{friend}(X, Z)}$$

is a shorthand for (given $U = \{tom, bob, sally, tina\}$)

$$\frac{\text{friend}(tom, bob) \wedge \text{friend}(bob, sally) : \text{friend}(tom, sally)}{\text{friend}(tom, sally)}$$

$$\frac{\text{friend}(tom, bob) \wedge \text{friend}(bob, tina) : \text{friend}(tom, tina)}{\text{friend}(tom, tina)}$$

...

Default theory (syntax of default logic)

Definition

A default theory T is a tuple $T = (W, \Delta)$ with

- ▶ $W \subseteq \mathcal{L}(\Sigma, V)$ (facts)
- ▶ Δ set of default rules

Main idea of default reasoning (semantics)

- ▶ apply defaults on W to extend the knowledge with plausible additional information
- ▶ apply defaults as long as no further information can be added

The resulting set of classical formulas is called *extension* and represents a plausible “belief state” of the default theory.

Main challenges

$$\delta = \frac{\phi : \psi_1, \dots, \psi_n}{\chi}$$

Meaning: If ϕ is known and ψ_1, \dots, ψ_n can be consistently assumed, then conclude χ .

When is such a default δ applicable, i. e.

- ▶ When is ϕ known?
- ▶ When can ψ_1, \dots, ψ_n be consistently assumed?

Approach: use a (at first unknown) extension E to check these conditions:

- ▶ ϕ is known iff $\phi \in E$;
- ▶ ψ_1, \dots, ψ_n can be consistently assumed iff $\neg\psi_i \notin E$,
 $1 \leq i \leq n$.

An extension $E \subseteq \mathcal{L}(\Sigma, V)$ is characterised by the following properties:

- ▶ E contains all facts: $W \subseteq E$
- ▶ E is deductively closed: $Cn(E) = E$
- ▶ E is *closed under default application*, i. e. if $\delta = \frac{\phi:\psi_1,\dots,\psi_n}{\chi} \in \Delta$ is applicable in E then $\chi \in E$ where:
 δ is applicable in E iff $\phi \in E$ and $\neg\psi_1 \notin E, \dots, \neg\psi_n \notin E$

More general:

- ▶ Let F be a deductively closed set of formulas
- ▶ Let K be a set of formulas (the context)

A default $\delta = \frac{\phi:\psi_1,\dots,\psi_n}{\chi}$ is *applicable* in F wrt. K iff

$$\phi \in F \text{ and } \neg\psi_1 \notin K, \dots, \neg\psi_n \notin K$$

Semantics of default logic: extensions 3/3

Let $T = (W, \Delta)$ be a default theory and S a set of formulas.

Define $\Lambda_T(S)$ to be the smallest set of formulas with

- ▶ $\Lambda_T(S)$ is deductively closed
- ▶ $W \subseteq \Lambda_T(S)$
- ▶ $\Lambda_T(S)$ is closed under default application wrt. the context S , i. e. for all $\delta = \frac{\phi:\psi_1,\dots,\psi_n}{\chi} \in \Delta$, if $\phi \in \Lambda_T(S)$ and $\neg\psi_1 \notin S, \dots, \neg\psi_n \notin S$, then $\chi \in \Lambda_T(S)$.

Definition

E is an *extension* of $T = (W, \Delta)$ iff $\Lambda_T(E) = E$.

Remark: conceptually, extensions of default theories correspond to models of a classical logic formula.

Example

$$T = (\{aquatic_creature\}, \{\frac{aquatic_creature : fish}{fish}\})$$

$E = Cn(\{aquatic_creature, fish\})$ is an extension of T ,

$E' = Cn(\{aquatic_creature, \neg fish\})$ is *not* an extension of T ,
despite the fact that

- ▶ $\{aquatic_creature\} \subseteq E'$
- ▶ E' is deductively closed
- ▶ E' is closed under default application

but $\Lambda_T(E') = Cn(\{aquatic_creature\}) \neq E'$

Example

$$T = (\{p\}, \{\frac{p:q}{q}, \frac{q:r}{r}\})$$

► $E = Cn(\{p, q, r\})$

Example

$$T = (\{p\}, \{\frac{p:q}{q}, \frac{p:\neg q}{\neg q}\})$$

► $E_1 = Cn(\{p, q\})$

► $E_2 = Cn(\{p, \neg q\})$

Example

$$T = (\{p\}, \{\frac{p:r}{s}\})$$

► $E = Cn(\{p, s\})$

A characterisation of extensions

Theorem (Reiter, 1980)

Let E be a set of closed formulas and let $T = (W, \Delta)$ be a default theory. Define a sequence of sets of formulas E_i , $i \geq 0$ via

- ▶ $E_0 = W$
- ▶ $E_{i+1} = Cn(E_i) \cup \{ \chi \mid \frac{\phi:\psi_1,\dots,\psi_n}{\chi} \in \Delta, \phi \in E_i, \neg\psi_1 \notin E, \dots, \neg\psi_n \notin E \}$

Then E is an extension of T iff

$$E = \bigcup_{i=0}^{\infty} E_i$$

Task: given $T = (W, \Delta)$, enumerate all extensions E of T

- ▶ We discuss an algorithm to explicitly enumerate all extensions using *default processes* and *process trees*
- ▶ Idea: apply defaults successively
 - ▶ as long as possible
 - ▶ use backtracking if inconsistency occurs

Let $T = (W, \Delta)$ be a fixed default theory.

Default sequences

$\Pi = (\delta_0, \dots, \delta_m)$ sequence of defaults from Δ (finite, no repetitions)

$\Pi[k] = (\delta_0, \dots, \delta_{k-1})$ sub-sequence of the first k elements

Definition

Let Π be a default sequence, define

$$\begin{aligned} In(\Pi) &= Cn(W \cup \{cons(\delta) \mid \delta \in \Pi\}) \\ Out(\Pi) &= \{\neg\psi \mid \psi \in just(\delta), \delta \in \Pi\} \end{aligned}$$

- ▶ $In(\Pi)$ collects formulas that are concluded by applying defaults from Π ; represents the current belief state after processing Π
- ▶ $Out(\Pi)$ collects formulas that should be not be proven true later

Default sequences - example

$T = (W, \Delta)$ with

$$W = \{a\} \quad \Delta = \{\delta_1 = \frac{a : \neg b}{\neg b}, \delta_2 = \frac{a : c}{c}\}$$

Let $\Pi_1 = (\delta_1)$, $\Pi_2 = (\delta_2, \delta_1)$.

$$In(\Pi_1) = Cn(\{a, \neg b\})$$

$$Out(\Pi_1) = \{b\}$$

$$In(\Pi_2) = Cn(\{a, c, \neg b\})$$

$$Out(\Pi_2) = \{\neg c, b\}$$

Note: $In(\Pi)$, $Out(\Pi)$ are independent of the actual order of defaults in Π

Definition

$\Pi = (\delta_0, \dots, \delta_m)$ is a *process* iff every δ_k is applicable in $In(\Pi[k])$.
In particular, δ_0 must be applicable in $In(()) = Cn(W)$.

A process Π is called

- ▶ *successful* iff $In(\Pi) \cap Out(\Pi) = \emptyset$;
- ▶ *not successful* (or *has failed*) iff $In(\Pi) \cap Out(\Pi) \neq \emptyset$;
- ▶ *closed* iff every $\delta \in \Delta$ that is applicable in $In(\Pi)$ appears in Π .

Theorem

E is an extension of T iff there is a closed and successful process Π with $E = In(\Pi)$.

Proof.

\Leftarrow :

Let Π be a closed and successful process of T with $E = \text{In}(\Pi)$.

We have to show $\Lambda_T(E) = E$:

► $\Lambda_T(E) \subseteq E$:

E is deductively closed and $W \subseteq E$; as Π is closed, E is closed wrt. default application. $\Lambda_T(E)$ is defined to be the *smallest* such set, so $\Lambda_T(E) \subseteq E$.

► $E \subseteq \Lambda_T(E)$:

By induction $\text{In}(\Pi[k]) \subseteq \Lambda_T(E)$ for all k . It follows $E \subseteq \Lambda_T(E)$.

\Rightarrow :

Let $E = \Lambda_T(E)$ be an extension of T .

Let $\Delta = \{\delta_0, \dots, \delta_n\}$ finite, arbitrary.

Construct process Π of T with $In(\Pi[k]) \subseteq E$ and $Out(\Pi[k]) \cap E = \emptyset$ as follows:

- ▶ $\Pi[0] = ()$;
- ▶ Let $\Pi[k]$ be constructed. If every default $\delta \in \Delta$ that is applicable in $In(\Pi[k])$ wrt. E is already in $\Pi[k]$, define $\Pi = \Pi[k]$ and halt. Otherwise select any default δ that is applicable in $In(\Pi[k])$ wrt. E and define $\Pi[k+1] = (\Pi[k], \delta)$.

In the end, $E = In(\Pi)$. □

Processes - examples

$T = (W, \Delta)$ with

$$W = \{a\}$$

$$\Delta = \{\delta_1 = \frac{a : \neg b}{\neg b}, \delta_2 = \frac{\top : c}{b}\}$$

- ▶ the process $\Pi_1 = (\delta_1)$
 - ▶ is successful: $In(\Pi_1) \cap Out(\Pi_1) = Cn(\{a, \neg b\}) \cap \{b\} = \emptyset$
 - ▶ is not closed as δ_2 is applicable in $In(\Pi_1)$
- ▶ the process $\Pi_2 = (\delta_1, \delta_2)$
 - ▶ is not successful:
 $In(\Pi_2) \cap Out(\Pi_2) = Cn(\{a, \neg b, b\}) \cap \{b, \neg c\} = \{b\} \neq \emptyset$
 - ▶ is closed
- ▶ the process $\Pi_3 = (\delta_2)$
 - ▶ is successful: $In(\Pi_3) \cap Out(\Pi_3) = Cn(\{a, b\}) \cap \{\neg c\} = \emptyset$
 - ▶ is closed
 - ▶ $E = In(\Pi_3) = Cn(\{a, b\})$ is extension of T

Process trees give an overview on all possible processes of a default theory $T = (W, \Delta)$:

- ▶ every node represents a process Π and is annotated with two labels: $In(\Pi)$ and $Out(\Pi)$;
- ▶ the root represents the empty process $\Pi = ()$ with $In(()) = Cn(W)$ and $Out(()) = \emptyset$;
- ▶ every application of a default induces a branch in the tree
- ▶ every leaf represents either
 - ▶ a failed process or
 - ▶ a closed and successful process

Construction of process trees


1. Construct a tree with a root N_0 and $In(N_0) = Cn(W)$ and $Out(N_0) = \emptyset$
2. As long as there is a leaf node N that is not marked with “failure” or “closed and successful”, repeat
 - ▶ If $In(N) \cap Out(N) \neq \emptyset$: mark node with “failure”
 - ▶ $In(N) \cap Out(N) = \emptyset$
 - ▶ for every applicable default $\delta = \frac{\phi:\psi_1,\dots,\psi_n}{\chi} \in \Delta$ that has not yet been considered in the process, add a child node $N(\delta)$ to N with

$$In(N(\delta)) = Cn(In(N) \cup \{\chi\})$$

$$Out(N(\delta)) = Out(N) \cup \{\neg\psi_1, \dots, \neg\psi_n\}$$

- ▶ Is there no further applicable default that has not yet been considered, mark N “closed and successful”; $In(N)$ is then an extension

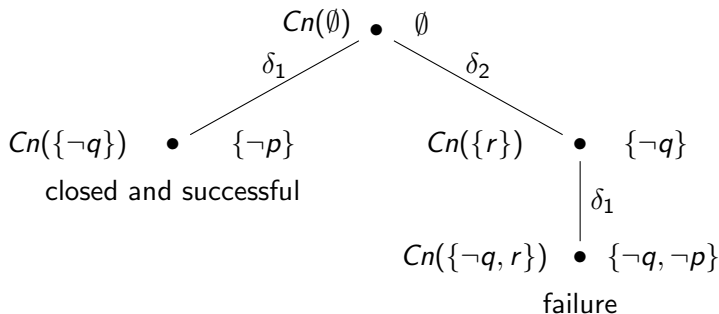
Process tree



$$\delta = \frac{\varphi : \psi_1, \dots, \psi_n}{\chi}$$


Process tree - example 1

$$T : \quad W = \emptyset \quad \Delta = \left\{ \delta_1 = \frac{\top : p}{\neg q}, \delta_2 = \frac{\top : q}{r} \right\}$$



Process tree - example 2

$$T : \quad W_0 = \{penguin \Rightarrow bird, penguin \Rightarrow \neg flies, bird\}$$
$$\Delta = \{\delta_1 = \frac{bird : flies}{flies}\}$$



For $W_1 = W_0 \cup \{penguin\}$
note:

$Cn(W_1)$ is the only extension
of $T_1 = (W_1, \Delta)$ as
 $\neg flies \in Cn(W_1)$.

Theorem (Minimality of extensions)

Let E, E' be extensions of a default theory T with $E \subseteq E'$. Then $E = E'$.

Theorem (Uniqueness of extensions)

Let $T = (W, \Delta)$ be a default theory a let

$$W \cup \{ \psi_1 \wedge \dots \wedge \psi_n \wedge \chi \mid \frac{\phi : \psi_1, \dots, \psi_n}{\chi} \in \Delta \}$$

be classically consistent. Then T has exactly one extension.

Theorem (Inconsistency 1)

A default theory $T = (W, \Delta)$ has an inconsistent extension iff W is already inconsistent.

Theorem (Inconsistency 2)

If T has an inconsistent extension E then E is the only extension of T .

$$T : \quad W = \emptyset \quad \Delta = \{\delta_0 = \frac{\top : a}{a}\}$$

T has exactly one extension $E = Cn(\{a\})$.

Add defaults to T :

- ▶ $\Delta_1 = \{\delta_0, \delta_1 = \frac{\top : b}{\neg b}\}$. $T_1 = (W, \Delta_1)$ has *no extensions*
- ▶ $\Delta_2 = \{\delta_0, \delta_2 = \frac{b : c}{c}\}$. $T_2 = (W, \Delta_2)$ has still exactly one extension E
- ▶ $\Delta_3 = \{\delta_0, \delta_3 = \frac{\top : \neg a}{\neg a}\}$. $T_3 = (W, \Delta_3)$ has two extensions E and $Cn(\{\neg a\})$.
- ▶ $\Delta_4 = \{\delta_0, \delta_4 = \frac{a : b}{b}\}$. $T_4 = (W, \Delta_4)$ has the extension $Cn(\{a, b\})$ which is a superset of E .

Extending a set of defaults (or the facts) can therefore

- ▶ remove extensions
- ▶ modify extensions
- ▶ create new extensions

Remember: this was the main motivation for default logic.

However, sometimes this behaviour may be too unpredictable.

Definition

Let $T = (W, \Delta)$ and $T' = (W, \Delta')$ default theories with the same set of facts and defaults $\Delta \subseteq \Delta'$. If every extension of T is contained in some extension of T' , then T' is called a *semi-monotone* extension of T .

In general, we can also not expect a semi-monotonic behaviour in default logic.

Definition

A default δ is called *normal* if $just(\delta) = cons(\delta)$, so δ is of the form

$$\delta = \frac{\phi : \psi}{\psi}$$

Example

$$\frac{bird : flies}{flies}$$

Using a normal default we can conclude a formula ψ if ψ is consistent with the current beliefs.

Processes of normal default theories

Let $T = (W, \Delta)$ be a normal default theory (=contains only normal defaults) with a consistent set of facts W .

Let $\Pi = (\delta_0, \dots, \delta_n)$ be a process of T with $\delta_i = \frac{\phi_i:\psi_i}{\psi_i}$.

Remember:

$$\begin{aligned} In(\Pi) &= Cn(W \cup \{\psi_i\}_{i \geq 0}) \\ Out(\Pi) &= \{\neg\psi_i\}_{i \geq 0} \end{aligned}$$

Can Π be a failure?

Every default δ_i was applicable, so $\neg\psi_i \notin In(\Pi)$. It follows $In(\Pi) \cap Out(\Pi) = \emptyset$ and therefore:

Theorem

Every process of a normal default theory is successful.

Theorem

A normal default theory always possesses at least one extension. Every finite process can be extended to a closed and successful process.

Theorem

Normal default theories are semi-monoton (adding another normal default extends previous extensions or preserves them completely).

Chapter 4.1: Default logic

Summary

- ▶ Default rules of the form

$$\delta = \frac{\phi : \psi_1, \dots, \psi_n}{\chi}$$

represent plausible (but not necessarily generally valid) rules

- ▶ Extensions: deductively closed, include facts, closed under default application
- ▶ Fix point characterisation of extensions
- ▶ Computing extensions with process trees
- ▶ Normal default theories and semi-monotony

- 1 Introduction
- 2 Classical logics and Prolog
- 3 Search and automatic planning
- 4 Knowledge representation and reasoning
 - Default logic
 - Answer set programming
 - Argumentation
- 5 Agents and multi agent systems
- 6 Summary and conclusion

- ▶ Default logic
 - ▶ ... is an expressive formalism for non-monotonic reasoning
 - ▶ ... is very formal and complex
 - ▶ therefore not very suitable for „practical knowledge representation“
- ▶ Recall Prolog
 - ▶ Practical programming language
 - ▶ Negation-as-failure `not`: similar to non-monotonic reasoning (if something cannot be proven, it is assumed to be false)
 - ▶ But there is no „logical negation“ in Prolog
 - ▶ Termination is not guaranteed
- ▶ Combine advantages of default logic and Prolog: *Answer set programming* (ASP)

Recall: some notation from first-order logic

- ▶ $\Sigma = (U, P, F)$ first-order signature, V set of variables
- ▶ In the following we consider only $F = \emptyset$
- ▶ A *literal* is an atom $p(t_1, \dots, t_k)$ or the negation of an atom $\neg p(t_1, \dots, t_k)$
- ▶ A literal ϕ is called *ground* if it mentions no variables

Extended logic programs: syntax 1/2

Let $\Sigma = (U, P, \emptyset)$ be a first-order signature and V a set of variables.

Definition

An *extended logic program* P is a (finite) set of rules of the form

$$r : \quad H \leftarrow A_1, \dots, A_n, \text{not } B_1, \dots, \text{not } B_m.$$

with literals $H, A_1, \dots, A_n, B_1, \dots, B_m$ from $\mathcal{L}(\Sigma, V)$.

- ▶ *not* is called *default negation*
- ▶ $\text{head}(r) = \{H\}$ is called head of the rule r
- ▶ $\text{pos}(r) = \{A_1, \dots, A_n\}$ positive body literals
- ▶ $\text{neg}(r) = \{B_1, \dots, B_m\}$ negative body literals

Negation-as-failure vs. classical negation

Why do we need two kinds of negation?

Compare

$cross_tracks \leftarrow not\ train_is_coming.$

“We can cross the tracks when we don’t know that a train is coming”

$cross_tracks \leftarrow \neg train_is_coming.$

“We can cross the tracks when we know that a train is not coming”
and

$call_doctor \leftarrow accident, not\ simulating.$

“We should call a doctor when there is an accident and we don’t know that the person is simulating his injury”

$call_doctor \leftarrow accident, \neg simulating.$

“We should call a doctor when there is an accident and we know that the person is not simulating his injury”

Extended logic programs: syntax 2/2

General rule

$$r : \quad H \leftarrow A_1, \dots, A_n, \text{not } B_1, \dots, \text{not } B_m.$$

Special cases:

- ▶ $n = m = 0$: Rule with empty body (=fact)

$$H \leftarrow . \quad \text{or simply} \quad H.$$

- ▶ All literals are atoms: normal logic rule
- ▶ Empty head literal ($\text{head}(r) = \emptyset$): constraint

$$\leftarrow A_1, \dots, A_n, \text{not } B_1, \dots, \text{not } B_m.$$

Example

$bird(X) \leftarrow penguin(X).$

$flies(X) \leftarrow bird(X), \text{not } \neg flies(X).$

$\neg flies(X) \leftarrow penguin(X).$

$flies(X) \leftarrow bat(X).$

$\leftarrow bird(X), bat(X).$

$penguin(tweety).$

$bat(batman).$

Grounding of extended logic programs

- ▶ An extended logic program with variables is always interpreted as a schema for its instances
- ▶ *Grounding* a program means substituting all variables by constants in all combinations

Example

For $U = \{a, b\}$ the program $P = \{p(X) \leftarrow t(X, Y), \text{not } r(Y), r(a).\}$ is a shorthand for $\text{ground}(P)$:

$$p(a) \leftarrow t(a, a), \text{not } r(a).$$
$$p(b) \leftarrow t(b, b), \text{not } r(b).$$
$$p(a) \leftarrow t(a, b), \text{not } r(b).$$
$$p(b) \leftarrow t(b, a), \text{not } r(a).$$
$$r(a).$$

→ it suffices to consider only propositional literals.

- ▶ Literals p and $\neg p$ are called *complementary*
- ▶ For a literal l , the complementary literal is denoted \bar{l}
- ▶ So $\bar{a} = \neg a$ and $\neg \bar{a} = a$ for an atom a
- ▶ A set of ground literals S is called *consistent* iff it contains no pair of complementary literals

Definition

A *state* is a consistent set of ground literals.

Conceptually, states of extended logic programs correspond to classical interpretations of classical formulas.

Example

Consider the extended logic program P :

$$p \leftarrow q, \text{not } r.$$
$$q \leftarrow \text{not } s.$$
$$s.$$

Some states for P are:

► $Z_1 = \{p, q, \neg r\}$

► $Z_2 = \{\neg p, s\}$

► $Z_3 = \{s, q, p\}$

Question: When does a state describe a program “in a meaningful way”?

Let P be an extended logic program *without* default negation.
Let S be a state.

Definition

S is *closed* under P iff for every rule $r \in P$, if $\text{pos}(r) \subseteq S$ then $\text{head}(r) \cap S \neq \emptyset$.

- ▶ For rules of the form $r : H \leftarrow A_1, \dots, A_n$ this means:
If $\{A_1, \dots, A_n\} \subseteq S$ then $H \in S$.
- ▶ For constraints of the form $r : \leftarrow A_1, \dots, A_n$ this means:
As $\text{head}(r) = \emptyset$ it has to hold $\text{head}(r) \cap S = \emptyset$; therefore $\{A_1, \dots, A_n\} \subseteq S$ must not be true.
- ▶ For facts $r : H \leftarrow$ this means:
As $\text{pos}(r) = \emptyset$ we always have $\text{pos}(r) \subseteq S$; every closed state must contain all facts.

Example

Consider the extended logic program P (without default negation):

$$p \leftarrow q, r.$$

$$q \leftarrow \neg s.$$

$$\neg s.$$

The following states are not closed:

- ▶ $Z_1 = \emptyset$
- ▶ $Z_2 = \{\neg s\}$
- ▶ $Z_3 = \{r, p\}$

The following states are closed:

- ▶ $Z_4 = \{\neg s, q\}$
- ▶ $Z_5 = \{\neg s, q, p\}$
- ▶ $Z_6 = \{\neg s, q, p, r\}$

Question: are Z_5 and Z_6 meaningful?

Minimal models

Let P be an extended logic program *without* default negation.
Let S be a state.

Definition

S is a *minimal model* of P iff S is closed and for every closed state S' for P , $S \subseteq S'$.

Example

Consider again P :

$$p \leftarrow q, r.$$

$$q \leftarrow \neg s.$$

$$\neg s.$$

Here $Z_4 = \{\neg s, q\}$ is the (only) minimal model of P .

Existence and uniqueness of minimal models 1/2

Let P be an extended logic program *without* default negation.

Definition

P is called *consistent* iff there is a closed state of P .

- ▶ $P_1 = \{s., \neg s.\}$ is not consistent (every closed state S would contain both s and $\neg s$; but a set containing complementary literals is not a state)
- ▶ $P_2 = \{s., \neg r \leftarrow s., r \leftarrow s.\}$ is not consistent.

Theorem

Let P be a consistent extended logic program without default negation. Then P has exactly one minimal model.

Existence and uniqueness of minimal models 2/2

Proof.

We have to show that there is at least one minimal model and at most one minimal model

- ▶ ≥ 1 : As P is consistent, there is a closed state S . If S is minimal: finished. If not, there is a another closed set $S' \subset S$. As P is finite there is a finite number of states and this sequence must end in a minimal model.
- ▶ ≤ 1 : Assume there are two different minimal models M_1, M_2 . Then $M_1 \not\subseteq M_2$ and $M_2 \not\subseteq M_1$ (otherwise one of them would not be minimal). We now show that $M_3 = M_1 \cap M_2$ is also closed:
 - ▶ Let $r \in P$ with $head(r) = H$ (analogous for constraints). If $pos(r) \subseteq M_3$ then $pos(r) \subseteq M_1$ and $pos(r) \subseteq M_2$. As M_1 and M_2 are closed, $H \in M_1$ and $H \in M_2$. Therefore $H \in M_3$.

As M_3 is closed, neither M_1 nor M_2 can be minimal (as $M_3 \subset M_1$ and $M_3 \subset M_2$). □

Characterisation of minimal models 1/5

Let P be an extended logic program *without* default negation.

Definition

For a set X of ground literals define

$$\Lambda_P(X) = \{ \text{head}(r) \mid r \in P, \text{pos}(r) \subseteq X \}$$

Example

Consider again $P = \{p \leftarrow q, r., q \leftarrow \neg s., \neg s.\}$. Then

- ▶ $\Lambda_P(\{r, q\}) = \{p, \neg s\}$
- ▶ $\Lambda_P(\{\neg s\}) = \{\neg s, q\}$
- ▶ $\Lambda_P(\emptyset) = \{\neg s\}$

Define also $\Lambda_P^1(X) = \Lambda_P(X)$ and $\Lambda_P^{n+1}(X) = \Lambda_P(\Lambda_P^n(X))$.

Theorem

Let P be a consistent extended logic program without default negation. A state S is closed under P iff $S \supseteq \Lambda_P(S)$.

Proof.

Let S be a state. If S is closed, there are no rule in P that can be applied. Therefore, all head literals H of all rules applicable in S are already in S . This is equivalent to $S \supseteq \Lambda_P(S)$. □

Theorem

Let P be a consistent extended logic program without default negation. Then there is a finite $k \geq 0$ with

$$\emptyset \subseteq \Lambda_P^1(\emptyset) \subseteq \Lambda_P^2(\emptyset) \subseteq \dots \subseteq \Lambda_P^k(\emptyset) = \Lambda_P^{k+1}(\emptyset) = \dots \quad (1)$$

and $\Lambda_P^k(\emptyset)$ is the minimal model of P .

Proof.

We first show (1):

We show this by induction.

- The base case $\emptyset \subseteq \Lambda_P^1(\emptyset)$ is obviously true.

- ▶ Assume $\emptyset \subseteq \Lambda_P^1(\emptyset) \subseteq \Lambda_P^2(\emptyset) \subseteq \dots \subseteq \Lambda_P^i(\emptyset)$. We have to show that $\Lambda_P^i(\emptyset) \subseteq \Lambda_P^{i+1}(\emptyset)$:

Let $H \in \Lambda_P^i(\emptyset)$. Then there is $l \leq i$ such that $H \in \Lambda_P^l(\emptyset)$ but $H \notin \Lambda_P^{l-1}(\emptyset)$. Let $r : H \leftarrow A_1, \dots, A_n \in P$ be a rule that caused $H \in \Lambda_P^l(\emptyset)$. Therefore $\{A_1, \dots, A_n\} \subseteq \Lambda_P^{l-1}(\emptyset)$. As $\Lambda_P^{l-1}(\emptyset) \subseteq \Lambda_P^i(\emptyset)$ we also have $\{A_1, \dots, A_n\} \in \Lambda_P^i(\emptyset)$. So we can apply r in calculating $\Lambda_P^{i+1}(\emptyset)$ and therefore $H \in \Lambda_P^{i+1}(\emptyset)$. Hence we get $\Lambda_P^i(\emptyset) \subseteq \Lambda_P^{i+1}(\emptyset)$.

It is also clear that this chain must end in a fixed point (the set of literals is finite), i. e., $\Lambda_P^k(\emptyset) = \Lambda_P^{k+1}(\emptyset)$.

We show now that $\Lambda_P^k(\emptyset)$ is a minimal model:

First, $M = \Lambda_P^k(\emptyset)$ is closed (as $M \supseteq \Lambda_P(M)$). Assume there is $M' \subset M$ such that M' is closed. Let $i \geq 0$ be the smallest index with

$$\begin{aligned}\Lambda_P^i(\emptyset) &\subseteq M' && \text{und} \\ \Lambda_P^{i+1}(\emptyset) &\not\subseteq M'\end{aligned}$$

Then there is $H \in \Lambda_P^{i+1}(\emptyset) \setminus M'$ such that there is a $r : H \leftarrow A_1, \dots, A_n \in P$ that was applicable when calculating $\Lambda_P^{i+1}(\emptyset)$, so $\{A_1, \dots, A_n\} \subseteq \Lambda_P^i(\emptyset)$. Therefore $\{A_1, \dots, A_n\} \subseteq M'$ as well and as $H \notin M'$, M' cannot be closed. □

... back to the general case of extended logic programs P with default negation.

- ▶ Idea: Simplify P to a program P' without default negation
- ▶ Compute the minimal model M of P' and call M *answer set* of P
- ▶ More specifically:
 1. „Guess“ a state S that could be an answer set
 2. Simplify P using S
 3. Compute the minimal model of the simpler program; if this turns out to be S again then S is an answer set

Definition

Let P be an extended logic program (with default negation) and S a state. The *reduct* P^S of P wrt. S is a logic program defined as

$$P^S = \{ H \leftarrow A_1, \dots, A_n \mid \\ H \leftarrow A_1, \dots, A_n, \text{not } B_1, \dots, \text{not } B_m. \in P, \\ \{B_1, \dots, B_m\} \cap S = \emptyset \}$$

The reduct P^S is constructed from P in two steps:

1. All rules that contain some `not` B with $B \in S$ in their body are removed.
2. For the remaining rules, all negative body literals are removed.

Observations:

- ▶ P^S looks different, depending on S
- ▶ P^S is an extended logic program without default negation
- ▶ P^S always contains
 - ▶ All facts from P
 - ▶ All rules without default negation

Remark: the reduct is also called the Gelfond-Lifschitz-Reduct after

Michael Gelfond, Vladimir Lifschitz. Classical negation in logic programs and disjunctive databases. In New Generation Computing 9:365–385, 1991.

Example

Consider the following program P

$$p \leftarrow \text{not } r.$$
$$r \leftarrow \neg q, \text{not } b.$$
$$\neg q \leftarrow b.$$
$$b.$$

and states S_1 and S_2

$$S_1 = \{r\}$$

$$S_2 = \{b, \neg q, p\}$$

Then

$$P^{S_1} = \{r \leftarrow \neg q. , \neg q \leftarrow b. , b.\}$$

$$P^{S_2} = \{p. , \neg q \leftarrow b. , b.\}$$

Definition

Let P be an extended logic program. A state S is an *answer set* of P iff S is the minimal model of P^S .

Example

As before:

$$P = \{p \leftarrow \text{not } r. \quad r \leftarrow \neg q, \text{not } b. \quad \neg q \leftarrow b. \quad b.\}$$

$$S_1 = \{r\}$$

$$S_2 = \{b, \neg q, p\}$$

$$P^{S_1} = \{r \leftarrow \neg q. \quad \neg q \leftarrow b. \quad b.\}$$

$$P^{S_2} = \{p. \quad \neg q \leftarrow b. \quad b.\}$$

- ▶ Minimal model of P^{S_1} is $\{b, \neg q, r\} \neq S_1$.
- ▶ Minimal model of P^{S_2} is $\{b, \neg q, p\} = S_2$, hence S_2 is an answer set of P .

Another example 1/3

Consider the extended logic program P :

$$p(X) \leftarrow \text{not } q(X).$$

$$q(X) \leftarrow r(X), \text{not } p(X).$$

$$r(a).$$

Let $U = \{a, b\}$. Grounding $P_g = \text{ground}(P)$ of P :

$$p(a) \leftarrow \text{not } q(a).$$

$$p(b) \leftarrow \text{not } q(b).$$

$$q(a) \leftarrow r(a), \text{not } p(a).$$

$$q(b) \leftarrow r(b), \text{not } p(b).$$

$$r(a).$$

Another example 2/3

$$\begin{aligned}P_g : & p(a) \leftarrow \text{not } q(a). \\ & p(b) \leftarrow \text{not } q(b). \\ & q(a) \leftarrow r(a), \text{not } p(a). \\ & q(b) \leftarrow r(b), \text{not } p(b). \\ & r(a).\end{aligned}$$

Assumption: $S_1 = \{r(a), q(a), p(b)\}$ is answer set. Compute reduct:

$$\begin{aligned}P_g^{S_1} : & p(b). \\ & q(a) \leftarrow r(a). \\ & r(a).\end{aligned}$$

Minimal model of $P_g^{S_1}$ is $S_1 \longrightarrow$ answer set.

Another example 3/3

$$\begin{aligned}P_g : & p(a) \leftarrow \text{not } q(a). \\ & p(b) \leftarrow \text{not } q(b). \\ & q(a) \leftarrow r(a), \text{not } p(a). \\ & q(b) \leftarrow r(b), \text{not } p(b). \\ & r(a).\end{aligned}$$

Assumption: $S_2 = \{r(a), p(a), p(b)\}$ is answer set. Compute reduct:

$$\begin{aligned}P_g^{S_2} : & p(a). \\ & p(b). \\ & r(a).\end{aligned}$$

Minimal model of $P_g^{S_2}$ is $S_2 \longrightarrow$ answer set.

Let P be an extended logic program.

For every rule $r : H \leftarrow A_1, \dots, A_n, \text{not } B_1, \dots, \text{not } B_m$
define the default

$$\text{def}(r) = \frac{A_1 \wedge \dots \wedge A_n : \overline{B_1}, \dots, \overline{B_m}}{H}$$

and $\text{def}(P) = (\emptyset, \{\text{def}(r) \mid r \in P\})$ as default theory wrt. P .

Theorem (Gelfond, Lifschitz, 1991)

If S is answer set of P then $\text{Cn}(S)$ is an extension of $\text{def}(P)$. If E is an extension of $\text{def}(P)$ then there is an answer set S von P with $E = \text{Cn}(S)$.

Chapter 4.2: Answer set programming

Summary

- ▶ extended logic programs contain rules of the form

$$r : \quad H \leftarrow A_1, \dots, A_n, \text{not } B_1, \dots, \text{not } B_m.$$

- ▶ grounding of first-order rules
- ▶ states, closed states
- ▶ minimal models of programs without default negation
- ▶ Gelfond-Lifschitz-Reduct and answer sets
- ▶ Answer sets and default extensions